

# Hidden Markov Models and Gaussian Mixture Models

Steve Renals and Peter Bell

Automatic Speech Recognition— ASR Lectures 4&5  
28/31 January 2013

## Overview

### HMMs and GMMs

- Key models and algorithms for HMM acoustic models
- Gaussians
- GMMs: Gaussian mixture models
- HMMs: Hidden Markov models
- HMM algorithms
  - Likelihood computation (forward algorithm)
  - Most probable state sequence (Viterbi algorithm)
  - Estimating the parameters (EM algorithm)

## Fundamental Equation of Statistical Speech Recognition

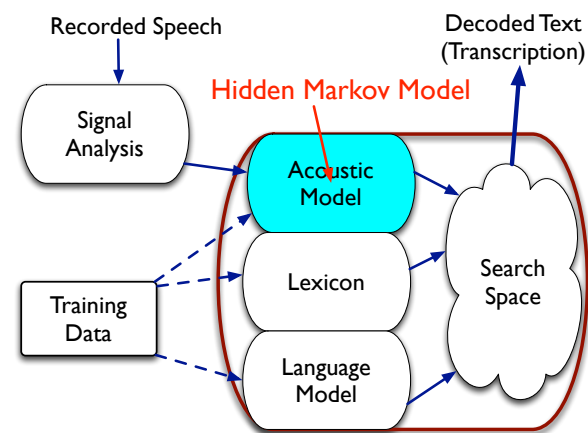
If  $\mathbf{X}$  is the sequence of acoustic feature vectors (observations) and  $\mathbf{W}$  denotes a word sequence, the most likely word sequence  $\mathbf{W}^*$  is given by

$$\mathbf{W}^* = \arg \max_{\mathbf{W}} P(\mathbf{W} | \mathbf{X})$$

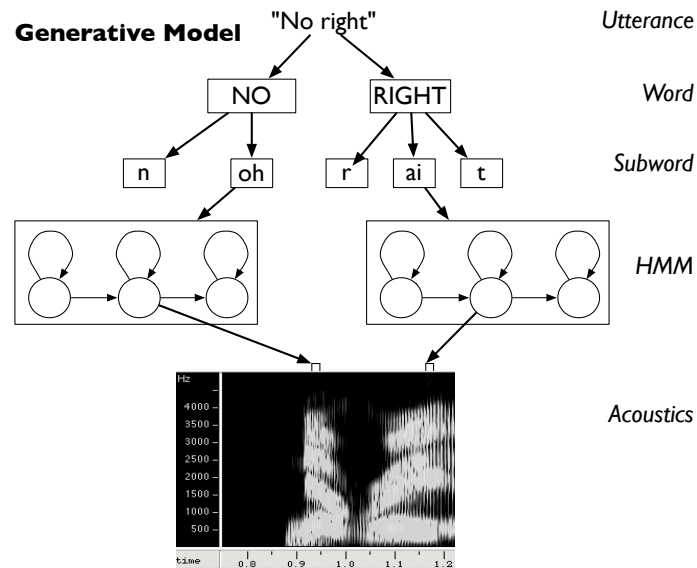
Applying Bayes' Theorem:

$$P(\mathbf{W} | \mathbf{X}) = \frac{p(\mathbf{X} | \mathbf{W})P(\mathbf{W})}{p(\mathbf{X})}$$
$$\propto p(\mathbf{X} | \mathbf{W})P(\mathbf{W})$$
$$\mathbf{W}^* = \arg \max_{\mathbf{W}} \underbrace{p(\mathbf{X} | \mathbf{W})}_{\text{Acoustic model}} \underbrace{P(\mathbf{W})}_{\text{Language model}}$$


## Acoustic Modelling



## Hierarchical modelling of speech



## Acoustic Model: Continuous Density HMM




Probabilistic finite state automaton

Parameters  $\lambda$ :

- Transition probabilities:  $a_{kj} = P(s_j | s_k)$
- Output probability density function:  $b_j(\mathbf{x}) = p(\mathbf{x} | s_j)$

## Acoustic Model: Continuous Density HMM




Probabilistic finite state automaton

Parameters  $\lambda$ :


- Transition probabilities:  $a_{kj} = P(s_j | s_k)$
- Output probability density function:  $b_j(\mathbf{x}) = p(\mathbf{x} | s_j)$

## HMM Assumptions



- 1 **Observation independence** An acoustic observation  $\mathbf{x}$  is conditionally independent of all other observations given the state that generated it
- 2 **Markov process** A state is conditionally independent of all other states given the previous state


## HMM Assumptions



- 1 **Observation independence** An acoustic observation  $\mathbf{x}$  is conditionally independent of all other observations given the state that generated it
- 2 **Markov process** A state is conditionally independent of all other states given the previous state

## HMM OUTPUT DISTRIBUTION

## Output distribution



Single multivariate Gaussian with mean  $\mu^j$ , covariance matrix  $\Sigma^j$ :

$$b_j(\mathbf{x}) = p(\mathbf{x} | s_j) = \mathcal{N}(\mathbf{x}; \mu^j, \Sigma^j)$$

M-component Gaussian mixture model:

$$b_j(\mathbf{x}) = p(\mathbf{x} | s_j) = \sum_{m=1}^M c_{jm} \mathcal{N}(\mathbf{x}; \mu^{jm}, \Sigma^{jm})$$

## Background: cdf

Consider a real valued random variable  $X$

- Cumulative distribution function (cdf)  $F(x)$  for  $X$ :

$$F(x) = P(X \leq x)$$

- To obtain the probability of falling in an interval we can do the following:

$$\begin{aligned} P(a < X \leq b) &= P(X \leq b) - P(X \leq a) \\ &= F(b) - F(a) \end{aligned}$$

- The rate of change of the cdf gives us the *probability density function* (pdf),  $p(x)$ :

$$p(x) = \frac{d}{dx}F(x) = F'(x)$$

$$F(x) = \int_{-\infty}^x p(x)dx$$

- $p(x)$  is **not** the probability that  $X$  has value  $x$ . But the pdf is proportional to the probability that  $X$  lies in a small interval centred on  $x$ .
- Notation:  $p$  for pdf,  $P$  for probability


- The **Gaussian** (or **Normal**) distribution is the most common (and easily analysed) continuous distribution
- It is also a reasonable model in many situations (the famous “bell curve”)
- If a (scalar) variable has a Gaussian distribution, then it has a probability density function with this form:

$$p(x|\mu, \sigma^2) = N(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

- The Gaussian is described by two parameters:
  - the mean  $\mu$  (location)
  - the variance  $\sigma^2$  (dispersion)


## Plot of Gaussian distribution

- Gaussians have the same shape, with the location controlled by the mean, and the spread controlled by the variance
- One-dimensional Gaussian with zero mean and unit variance ( $\mu = 0, \sigma^2 = 1$ ):



## Properties of the Gaussian distribution

$$N(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$



## Parameter estimation

- Estimate mean and variance parameters of a Gaussian from data  $x^1, x^2, \dots, x^n$
- Use sample mean and sample variance estimates:

$$\mu = \frac{1}{n} \sum_{i=1}^n x^i \quad (\text{sample mean})$$

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x^i - \mu)^2 \quad (\text{sample variance})$$

## Exercise

Consider the log likelihood of a set of  $N$  data points  $\{x^1, \dots, x^N\}$  being generated by a Gaussian with mean  $\mu$  and variance  $\sigma^2$ :

$$\begin{aligned} L = \ln p(\{x^1, \dots, x^N\} | \mu, \sigma^2) &= -\frac{1}{2} \sum_{n=1}^N \left( \frac{(x_n - \mu)^2}{\sigma^2} - \ln \sigma^2 - \ln(2\pi) \right) \\ &= -\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi) \end{aligned}$$

By maximising the the log likelihood function with respect to  $\mu$  show that the maximum likelihood estimate for the mean is indeed the sample mean:

$$\mu_{ML} = \frac{1}{N} \sum_{n=1}^N x_n.$$

## The multidimensional Gaussian distribution

- The  $d$ -dimensional vector  $\mathbf{x}$  is multivariate Gaussian if it has a probability density function of the following form:

$$p(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \exp \left( -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right)$$

The pdf is parameterized by the mean vector  $\boldsymbol{\mu}$  and the covariance matrix  $\boldsymbol{\Sigma}$ .

- The 1-dimensional Gaussian is a special case of this pdf
- The argument to the exponential  $0.5(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})$  is referred to as a *quadratic form*.

## Covariance matrix

- The mean vector  $\boldsymbol{\mu}$  is the expectation of  $\mathbf{x}$ :

$$\boldsymbol{\mu} = E[\mathbf{x}]$$

- The covariance matrix  $\boldsymbol{\Sigma}$  is the expectation of the deviation of  $\mathbf{x}$  from the mean:


$$\boldsymbol{\Sigma} = E[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T]$$

- $\boldsymbol{\Sigma}$  is a  $d \times d$  symmetric matrix:

$$\Sigma_{ij} = E[(x_i - \mu_i)(x_j - \mu_j)] = E[(x_j - \mu_j)(x_i - \mu_i)] = \Sigma_{ji}$$


- The sign of the covariance helps to determine the relationship between two components:
  - If  $x_j$  is large when  $x_i$  is large, then  $(x_j - \mu_j)(x_i - \mu_i)$  will tend to be positive;
  - If  $x_j$  is small when  $x_i$  is large, then  $(x_j - \mu_j)(x_i - \mu_i)$  will tend to be negative.

## Spherical Gaussian




$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \quad \rho_{12} = 0$$

## Diagonal Covariance Gaussian



$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 4 \end{pmatrix} \quad \rho_{12} = 0$$

## Full covariance Gaussian



$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 1 & -1 \\ -1 & 4 \end{pmatrix} \quad \rho_{12} = -0.5$$

## Parameter estimation


- It is possible to show that the mean vector  $\hat{\mu}$  and covariance matrix  $\hat{\Sigma}$  that maximize the likelihood of the training data are given by:

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}^n$$


$$\hat{\Sigma} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}^n - \hat{\mu})(\mathbf{x}^n - \hat{\mu})^T$$

- The mean of the distribution is estimated by the sample mean and the covariance by the sample covariance


## Example data



## Maximum likelihood fit to a Gaussian




## Data in clusters (example 1)



$$\mu_1 = [0 \ 0]^T \quad \mu_2 = [1 \ 1]^T \quad \Sigma_1 = \Sigma_2 = 0.2I$$

## Example 1 fit by a Gaussian




$$\mu_1 = [0 \ 0]^T \quad \mu_2 = [1 \ 1]^T \quad \Sigma_1 = \Sigma_2 = 0.2I$$


## k-means clustering

- k-means is an automatic procedure for clustering unlabelled data
- Requires a prespecified number of clusters
- Clustering algorithm chooses a set of clusters with the minimum within-cluster variance
- Guaranteed to converge (eventually)
- Clustering solution is dependent on the initialisation


## k-means example: data set



## k-means example: initialization




## k-means example: iteration 1 (assign points to clusters)







## k-means example: iteration 1 (recompute centres)




## k-means example: iteration 2 (assign points to clusters)



## k-means example: iteration 2 (recompute centres)



## k-means example: iteration 3 (assign points to clusters)



No changes, so converged

## Mixture model

- A more flexible form of density estimation is made up of a linear combination of component densities:

$$p(\mathbf{x}) = \sum_{j=1}^M p(\mathbf{x}|j)P(j)$$

- This is called a *mixture model* or a *mixture density*
- $p(\mathbf{x}|j)$ : component densities
- $P(j)$ : mixing parameters
- Generative model:
  - Choose a mixture component based on  $P(j)$
  - Generate a data point  $\mathbf{x}$  from the chosen component using  $p(\mathbf{x}|j)$

## Component occupation probability

- We can apply Bayes' theorem:

$$P(j|\mathbf{x}) = \frac{p(\mathbf{x}|j)P(j)}{p(\mathbf{x})} = \frac{p(\mathbf{x}|j)P(j)}{\sum_{j=1}^M p(\mathbf{x}|j)P(j)}$$

- The posterior probabilities  $P(j|\mathbf{x})$  give the probability that component  $j$  was responsible for generating data point  $\mathbf{x}$
- The  $P(j|\mathbf{x})$ s are called the *component occupation probabilities* (or sometimes called the *responsibilities*)
- Since they are posterior probabilities:

$$\sum_{j=1}^M P(j|\mathbf{x}) = 1$$


## Parameter estimation

- If we knew which mixture component was responsible for a data point:
  - we would be able to assign each point unambiguously to a mixture component
  - and we could estimate the mean for each component Gaussian as the sample mean (just like k-means clustering)
  - and we could estimate the covariance as the sample covariance
- But we don't know which mixture component a data point comes from...
- Maybe we could use the component occupation probabilities  $P(j|\mathbf{x})$ ?

## Gaussian mixture model

- The most important mixture model is the *Gaussian Mixture Model* (GMM), where the component densities are Gaussians
- Consider a GMM, where each component Gaussian  $N_j(\mathbf{x}; \mu_j, \sigma_j^2)$  has mean  $\mu_j$  and a spherical covariance  $\Sigma = \sigma^2 \mathbf{I}$

$$p(\mathbf{x}) = \sum_{j=1}^P P(j)p(\mathbf{x}|j) = \sum_{j=1}^P P(j)N_j(\mathbf{x}; \mu_j, \sigma_j^2)$$



## GMM Parameter estimation when we know which component generated the data

- Define the indicator variable  $z_{jn} = 1$  if component  $j$  generated component  $\mathbf{x}^n$  (and 0 otherwise)
- If  $z_{jn}$  wasn't hidden then we could count the number of observed data points generated by  $j$ :

$$N_j = \sum_{n=1}^N z_{jn}$$

- And estimate the mean, variance and mixing parameters as:

$$\hat{\mu}_j = \frac{\sum_n z_{jn} \mathbf{x}^n}{N_j}$$

$$\hat{\sigma}_j^2 = \frac{\sum_n z_{jn} \|\mathbf{x}^n - \mu_k\|^2}{N_j}$$

$$\hat{P}(j) = \frac{1}{N} \sum_n z_{jn} = \frac{N_j}{N}$$

## Soft assignment

- Estimate "soft counts" based on the component occupation probabilities  $P(j|\mathbf{x}^n)$ :

$$N_j^* = \sum_{n=1}^N P(j|\mathbf{x}^n)$$

- We can imagine assigning data points to component  $j$  weighted by the component occupation probability  $P(j|\mathbf{x}^n)$
- So we could imagine estimating the mean, variance and prior probabilities as:

$$\hat{\mu}_j = \frac{\sum_n P(j|\mathbf{x}^n) \mathbf{x}^n}{\sum_n P(j|\mathbf{x}^n)} = \frac{\sum_n P(j|\mathbf{x}^n) \mathbf{x}^n}{N_j^*}$$

$$\hat{\sigma}_j^2 = \frac{\sum_n P(j|\mathbf{x}^n) \|\mathbf{x}^n - \mu_k\|^2}{\sum_n P(j|\mathbf{x}^n)} = \frac{\sum_n P(j|\mathbf{x}^n) \|\mathbf{x}^n - \mu_k\|^2}{N_j^*}$$

$$\hat{P}(j) = \frac{1}{N} \sum_n P(j|\mathbf{x}^n) = \frac{N_j^*}{N}$$

## EM algorithm

- Problem!* Recall that:

$$P(j|\mathbf{x}) = \frac{p(\mathbf{x}|j)P(j)}{p(\mathbf{x})}$$

We need to know  $p(\mathbf{x}|j)$  and  $P(j)$  to estimate the parameters of  $p(\mathbf{x}|j)$  and to estimate  $P(j)$ ....

- Solution: an iterative algorithm where each iteration has two parts:
  - Compute the component occupation probabilities  $P(j|\mathbf{x})$  using the current estimates of the GMM parameters (means, variances, mixing parameters) (E-step)
  - Computer the GMM parameters using the current estimates of the component occupation probabilities (M-step)
- Starting from some initialization (e.g. using k-means for the means) these steps are alternated until convergence
- This is called the *EM Algorithm* and can be shown to maximize the likelihood

## Maximum likelihood parameter estimation

- The likelihood of a data set  $\mathbf{X} = \{\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^N\}$  is given by:


$$\mathcal{L} = \prod_{n=1}^N p(\mathbf{x}^n) = \prod_{n=1}^N \sum_{j=1}^M p(\mathbf{x}^n|j)P(j)$$

- We can regard the *negative log likelihood* as an error function:


$$\begin{aligned} E &= -\ln \mathcal{L} = -\sum_{n=1}^N \ln p(\mathbf{x}^n) \\ &= -\sum_{n=1}^N \ln \left( \sum_{j=1}^M p(\mathbf{x}^n|j)P(j) \right) \end{aligned}$$

- Considering the derivatives of  $E$  with respect to the parameters, gives expressions like the previous slide

## Example 1 fit using a GMM




## Peakily distributed data (Example 2)





$$\mu_1 = \mu_2 = [0 \ 0]^T \quad \Sigma_1 = 0.1\mathbf{I} \quad \Sigma_2 = 2\mathbf{I}$$

## Example 2 fit by a Gaussian




$$\mu_1 = \mu_2 = [0 \ 0]^T \quad \Sigma_1 = 0.1\mathbf{I} \quad \Sigma_2 = 2\mathbf{I}$$

## Example 2 fit by a GMM




## Example 2: component Gaussians



## Comments on GMMs

- GMMs trained using the EM algorithm are able to self organize to fit a data set
- Individual components take responsibility for parts of the data set (probabilistically)
- Soft assignment to components not hard assignment — “soft clustering”
- GMMs scale very well, e.g.: large speech recognition systems can have 30,000 GMMs, each with 32 components: sometimes 1 million Gaussian components!! And the parameters all estimated from (a lot of) data by EM

## Back to HMMs...



Output distribution:

- Single multivariate Gaussian with mean  $\mu^j$ , covariance matrix  $\Sigma^j$ :

$$b_j(\mathbf{x}) = p(\mathbf{x} | s_j) = \mathcal{N}(\mathbf{x}; \mu^j, \Sigma^j)$$

- $M$ -component Gaussian mixture model:

$$b_j(\mathbf{x}) = p(\mathbf{x} | s_j) = \sum_{m=1}^M c_{jm} \mathcal{N}(\mathbf{x}; \mu^{jm}, \Sigma^{jm})$$

## The three problems of HMMs

Working with HMMs requires the solution of three problems:


- 1 **Likelihood** Determine the overall likelihood of an observation sequence  $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_t, \dots, \mathbf{x}_T)$  being generated by an HMM
- 2 **Decoding** Given an observation sequence and an HMM, determine the most probable hidden state sequence
- 3 **Training** Given an observation sequence and an HMM, learn the best HMM parameters  $\lambda = \{\{a_{jk}\}, \{b_j(\cdot)\}\}$

## 1. Likelihood: The Forward algorithm

- Goal: determine  $p(\mathbf{X} | \lambda)$
- Sum over all possible state sequences  $s_1 s_2 \dots s_T$  that could result in the observation sequence  $X$
- Rather than enumerating each sequence, compute the probabilities recursively (exploiting the Markov assumption)

## Recursive algorithms on HMMs

Visualize the problem as a *state-time trellis*



## 1. Likelihood: The Forward algorithm

- Goal: determine  $p(\mathbf{X} | \lambda)$
- Sum over all possible state sequences  $s_1 s_2 \dots s_T$  that could result in the observation sequence  $X$
- Rather than enumerating each sequence, compute the probabilities recursively (exploiting the Markov assumption)
- *Forward probability*,  $\alpha_t(s_j)$ : the probability of observing the observation sequence  $\mathbf{x}_1 \dots \mathbf{x}_t$  and being in state  $s_j$  at time  $t$ :

$$\alpha_t(s_j) = p(\mathbf{x}_1, \dots, \mathbf{x}_t, S(t) = s_j | \lambda)$$

## 1. Likelihood: The Forward recursion

- Initialization

$$\begin{aligned} \alpha_0(s_I) &= 1 \\ \alpha_0(s_j) &= 0 \quad \text{if } s_j \neq s_I \end{aligned}$$

- Recursion


$$\alpha_t(s_j) = \sum_{i=1}^N \alpha_{t-1}(s_i) a_{ij} b_j(\mathbf{x}_t)$$

- Termination

$$p(\mathbf{X} | \lambda) = \alpha_T(s_E) = \sum_{i=1}^N \alpha_T(s_i) a_{iE}$$

## 1. Likelihood: Forward Recursion

$$\alpha_t(s_j) = p(\mathbf{x}_1, \dots, \mathbf{x}_t, S(t) = s_j \mid \lambda)$$



## Viterbi approximation


- Instead of summing over all possible state sequences, just consider the most likely
- Achieve this by changing the summation to a maximisation in the recursion:

$$V_t(s_j) = \max_i V_{t-1}(s_i) a_{ij} b_j(\mathbf{x}_t)$$

- Changing the recursion in this way gives the likelihood of the most probable path
- We need to keep track of the states that make up this path by keeping a sequence of *backpointers* to enable a Viterbi *backtrace*: the backpointer for each state at each time indicates the previous state on the most probable path


## Viterbi Recursion

Likelihood of the most probable path



## Viterbi Recursion

Backpointers to the previous state on the most probable path



## 2. Decoding: The Viterbi algorithm

- Initialization

$$\begin{aligned} V_0(s_i) &= 1 \\ V_0(s_j) &= 0 \quad \text{if } s_j \neq s_i \\ bt_0(s_j) &= 0 \end{aligned}$$

- Recursion


$$\begin{aligned} V_t(s_j) &= \max_{i=1}^N V_{t-1}(s_i) a_{ij} b_j(\mathbf{x}_t) \\ bt_t(s_j) &= \arg \max_{i=1}^N V_{t-1}(s_i) a_{ij} b_j(\mathbf{x}_t) \end{aligned}$$

- Termination

$$\begin{aligned} P^* &= V_T(s_E) = \max_{i=1}^N V_T(s_i) a_{iE} \\ s_T^* &= bt_T(s_E) = \arg \max_{i=1}^N V_T(s_i) a_{iE} \end{aligned}$$

## Viterbi Backtrace

Backtrace to find the state sequence of the most probable path



## 3. Training: Forward-Backward algorithm

- Goal: Efficiently estimate the parameters of an HMM  $\lambda$  from an observation sequence
- Assume single Gaussian output probability distribution

$$b_j(\mathbf{x}) = p(\mathbf{x} | s_j) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}^j, \boldsymbol{\Sigma}^j)$$

- Parameters  $\lambda$ :

- Transition probabilities  $a_{ij}$ :

$$\sum_j a_{ij} = 1$$

- Gaussian parameters for state  $s_j$ :  
mean vector  $\boldsymbol{\mu}^j$ ; covariance matrix  $\boldsymbol{\Sigma}^j$

## Viterbi Training

- If we knew the state-time alignment, then each observation feature vector could be assigned to a specific state
- A state-time alignment can be obtained using the most probable path obtained by Viterbi decoding
- Maximum likelihood estimate of  $a_{ij}$ , if  $C(s_i \rightarrow s_j)$  is the count of transitions from  $s_i$  to  $s_j$

$$\hat{a}_{ij} = \frac{C(s_i \rightarrow s_j)}{\sum_k C(s_i \rightarrow s_k)}$$

- Likewise if  $Z_j$  is the set of observed acoustic feature vectors assigned to state  $j$ , we can use the standard maximum likelihood estimates for the mean and the covariance:

$$\begin{aligned} \hat{\boldsymbol{\mu}}^j &= \frac{\sum_{\mathbf{x} \in Z_j} \mathbf{x}}{|Z_j|} \\ \hat{\boldsymbol{\Sigma}}^j &= \frac{\sum_{\mathbf{x} \in Z_j} (\mathbf{x} - \hat{\boldsymbol{\mu}}^j)(\mathbf{x} - \hat{\boldsymbol{\mu}}^j)^T}{|Z_j|} \end{aligned}$$



## EM Algorithm

- Viterbi training is an approximation—we would like to consider *all* possible paths
- In this case rather than having a hard state-time alignment we estimate a probability
- *State occupation probability*: The probability  $\gamma_t(s_j)$  of occupying state  $s_j$  at time  $t$  given the sequence of observations.  
Compare with component occupation probability in a GMM
- We can use this for an iterative algorithm for HMM training: the EM algorithm
- Each iteration has two steps:
  - E-step** estimate the state occupation probabilities (Expectation)
  - M-step** re-estimate the HMM parameters based on the estimated state occupation probabilities (Maximisation)

## Backward probabilities

- To estimate the state occupation probabilities it is useful to define (recursively) another set of probabilities—the *Backward* probabilities

$$\beta_t(s_j) = p(\mathbf{x}_{t+1}, \mathbf{x}_{t+2}, \mathbf{x}_T \mid S(t) = s_j, \lambda)$$

The probability of future observations given a the HMM is in state  $s_j$  at time  $t$

- These can be recursively computed (going backwards in time)
  - Initialisation

$$\beta_T(s_i) = a_{iE}$$

- Recursion


$$\beta_t(s_i) = \sum_{j=1}^N a_{ij} b_j(\mathbf{x}_{t+1}) \beta_{t+1}(s_j)$$

- Termination

$$p(\mathbf{X} \mid \lambda) = \beta_0(s_i) = \sum_{j=1}^N a_{ij} b_j(\mathbf{x}_1) \beta_1(s_j) = \alpha_T(s_E)$$

## Backward Recursion

$$\beta_t(s_j) = p(\mathbf{x}_{t+1}, \mathbf{x}_{t+2}, \mathbf{x}_T \mid S(t) = s_j, \lambda)$$



## State Occupation Probability

- The **state occupation probability**  $\gamma_t(s_j)$  is the probability of occupying state  $s_j$  at time  $t$  given the sequence of observations
- Express in terms of the forward and backward probabilities:

$$\gamma_t(s_j) = P(S(t) = s_j \mid \mathbf{X}, \lambda) = \frac{1}{\alpha_T(s_E)} \alpha_t(j) \beta_t(j)$$

recalling that  $p(\mathbf{X} \mid \lambda) = \alpha_T(s_E)$

- Since

$$\begin{aligned} \alpha_t(s_j) \beta_t(s_j) &= p(\mathbf{x}_1, \dots, \mathbf{x}_t, S(t) = s_j \mid \lambda) \\ &= p(\mathbf{x}_1, \dots, \mathbf{x}_t, \mathbf{x}_{t+1}, \mathbf{x}_{t+2}, \dots, \mathbf{x}_T, S(t) = s_j \mid \lambda) \\ &= p(\mathbf{X}, S(t) = s_j \mid \lambda) \end{aligned}$$

$$P(S(t) = s_j \mid \mathbf{X}, \lambda) = \frac{p(\mathbf{X}, S(t) = s_j \mid \lambda)}{p(\mathbf{X} \mid \lambda)}$$

## Re-estimation of Gaussian parameters

- The sum of state occupation probabilities through time for a state, may be regarded as a “soft” count
- We can use this “soft” alignment to re-estimate the HMM parameters:

$$\hat{\mu}^j = \frac{\sum_{t=1}^T \gamma_t(s_j) \mathbf{x}_t}{\sum_{t=1}^T \gamma_t(s_j)}$$

$$\hat{\Sigma}^j = \frac{\sum_{t=1}^T \gamma_t(s_j) (\mathbf{x}_t - \hat{\mu}^j)(\mathbf{x}_t - \hat{\mu}^j)^T}{\sum_{t=1}^T \gamma_t(s_j)}$$

## Re-estimation of transition probabilities

- Similarly to the state occupation probability, we can estimate  $\xi_t(s_i, s_j)$ , the probability of being in  $s_i$  at time  $t$  and  $s_j$  at  $t + 1$ , given the observations:

$$\begin{aligned} \xi_t(s_i, s_j) &= P(S(t) = s_i, S(t+1) = s_j \mid \mathbf{X}, \lambda) \\ &= \frac{P(S(t) = s_i, S(t+1) = s_j, \mathbf{X} \mid \lambda)}{p(\mathbf{X} \mid \lambda)} \\ &= \frac{\alpha_t(s_i) a_{ij} b_j(\mathbf{x}_{t+1}) \beta_{t+1}(s_j)}{\alpha_T(s_E)} \end{aligned}$$

- We can use this to re-estimate the transition probabilities

$$\hat{a}_{ij} = \frac{\sum_{t=1}^T \xi_t(s_i, s_j)}{\sum_{k=1}^N \sum_{t=1}^T \xi_t(s_i, s_k)}$$

## Pulling it all together

- Iterative estimation of HMM parameters using the EM algorithm. At each iteration
  - E step** For all time-state pairs
    - Recursively compute the forward probabilities  $\alpha_t(s_j)$  and backward probabilities  $\beta_t(j)$
    - Compute the state occupation probabilities  $\gamma_t(s_j)$  and  $\xi_t(s_i, s_j)$
  - M step** Based on the estimated state occupation probabilities re-estimate the HMM parameters: mean vectors  $\mu^j$ , covariance matrices  $\Sigma^j$  and transition probabilities  $a_{ij}$
- The application of the EM algorithm to HMM training is sometimes called the Forward-Backward algorithm

## Extension to a corpus of utterances

- We usually train from a large corpus of  $R$  utterances
- If  $\mathbf{x}_t^r$  is the  $t$ th frame of the  $r$ th utterance  $\mathbf{X}^r$  then we can compute the probabilities  $\alpha_t^r(j)$ ,  $\beta_t^r(j)$ ,  $\gamma_t^r(s_j)$  and  $\xi_t^r(s_i, s_j)$  as before
- The re-estimates are as before, except we must sum over the  $R$  utterances, eg:

$$\hat{\mu}^j = \frac{\sum_{r=1}^R \sum_{t=1}^T \gamma_t^r(s_j) \mathbf{x}_t^r}{\sum_{r=1}^R \sum_{t=1}^T \gamma_t^r(s_j)}$$

## Extension to Gaussian mixture model (GMM)

- The assumption of a Gaussian distribution at each state is very strong; in practice the acoustic feature vectors associated with a state may be strongly non-Gaussian
- In this case an  $M$ -component Gaussian mixture model is an appropriate density function:

$$b_j(\mathbf{x}) = p(\mathbf{x} | s_j) = \sum_{m=1}^M c_{jm} \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}^{jm}, \boldsymbol{\Sigma}^{jm})$$

Given enough components, this family of functions can model any distribution.

- Train using the EM algorithm, in which the component estimation probabilities are estimated in the E-step

## EM training of HMM/GMM

- Rather than estimating the state-time alignment, we estimate the component/state-time alignment, and component-state occupation probabilities  $\gamma_t(s_j, m)$ : the probability of occupying mixture component  $m$  of state  $s_j$  at time  $t$
- We can thus re-estimate the mean of mixture component  $m$  of state  $s_j$  as follows

$$\hat{\boldsymbol{\mu}}^{jm} = \frac{\sum_{t=1}^T \gamma_t(s_j, m) \mathbf{x}_t}{\sum_{t=1}^T \gamma_t(s_j, m)}$$

And likewise for the covariance matrices (mixture models often use diagonal covariance matrices)

- The mixture coefficients are re-estimated in a similar way to transition probabilities:

$$\hat{c}_{jm} = \frac{\sum_{t=1}^T \gamma_t(s_j, m)}{\sum_{\ell=1}^M \sum_{t=1}^T \gamma_t(s_j, \ell)}$$

## Doing the computation

- The forward, backward and Viterbi recursions result in a long sequence of probabilities being multiplied
- This can cause floating point *underflow* problems
- In practice computations are performed in the log domain (in which multiplies become adds)
- Working in the log domain also avoids needing to perform the exponentiation when computing Gaussians

## Summary: HMMs

- HMMs provide a generative model for statistical speech recognition
- Three key problems
  - ① Computing the overall likelihood: the Forward algorithm
  - ② Decoding the most likely state sequence: the Viterbi algorithm
  - ③ Estimating the most likely parameters: the EM (Forward-Backward) algorithm
- Solutions to these problems are tractable due to the two key HMM assumptions
  - ① Conditional independence of observations given the current state
  - ② Markov assumption on the states

- Gales and Young (2007). “The Application of Hidden Markov Models in Speech Recognition”, *Foundations and Trends in Signal Processing*, **1** (3), 195–304: section 2.2.
- Jurafsky and Martin (2008). *Speech and Language Processing* (2nd ed.): sections 6.1–6.5; 9.2; 9.4. (Errata at <http://www.cs.colorado.edu/~martin/SLP/Errata/SLP2-PIEV-Errata.html>)
- Rabiner and Juang (1989). “An introduction to hidden Markov models”, *IEEE ASSP Magazine*, **3** (1), 4–16.
- Renals and Hain (2010). “Speech Recognition”, *Computational Linguistics and Natural Language Processing Handbook*, Clark, Fox and Lappin (eds.), Blackwells.